

WrightSim

Kyle Sunden

Goals

Theory

NISE

Algorithmic
Improvements

Parallel
Implementations

Scaling Analysis

Limitations



WrightSim

Kyle Sunden

University of Wisconsin–Madison

January 3, 2018

Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

- ▶ Reproduce experimental spectra *in silico*.
- ▶ Designed with experimentalists in mind.
- ▶ Uses numerical integration for flexibility, accuracy and interpretability.
- ▶ Focus on Frequency-domain spectroscopy, but techniques in principle extend to time-domain.
- ▶ Output retains frequency and phase information, can be combined with other simulations and measured similar to a monochromator.
- ▶ Selectivity in what portions of the overall signal are simulated, providing deeper understanding.



Goals


Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

Presented here is a description of *what* is done to perform these simulations, to understand *why* it works, please refer to Kohler, Thompson, and Wright**Kohler_2017**.

The simulation uses a set number of electric fields (3, in the case of the simulation presented here).

These electric fields interact in combinatorically large different fashions.

Interactions create superposition coherences and/or populations in material systems.



Goals


Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis


Limitations



There are six time orderings for interactions to occur (I-VI).

There are 16 independant pathways possible for two positive (solid up/dashed down) and one negative (dashed up/solid down) interactions. Originally from **Kohler_2017**.





Density elements, encoded with quantum state (subscript) and the electric fields which have interacted (superscript). Colored arrows represent the different electric fields. All possible states which have the desired conditions for the process simulated are included. These form the state vector (right). **Kohler_2017**



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

$$\overline{\overline{Q}} \equiv \begin{bmatrix} 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ -A_2 & -\Gamma_{10} & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ A_{2'} & 0 & -\Gamma_{10} & 0 & 0 & 0 & 0 & 0 & 0 \\ A_1 & 0 & 0 & -\Gamma_{10} & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & B_1 & B_{2'} & -\Gamma_{20} & 0 & 0 & 0 & 0 \\ 0 & A_1 & 0 & -A_2 & 0 & -\Gamma_{11} & 0 & 0 & 0 \\ 0 & A_{2'} & -A_2 & 0 & 0 & 0 & -\Gamma_{11} & 0 & 0 \\ 0 & 0 & 0 & 0 & B_2 & -2A_{2'} & -2A_1 & -\Gamma_{10} & 0 \\ 0 & 0 & 0 & 0 & -A_2 & B_{2'} & B_1 & 0 & -\Gamma_{21} \end{bmatrix}$$

Defines the transition between states, dependant on the electric field. Γ represents the dephasing/population decay. A and B variables incorporate the dipole moment and electric field terms. **Kohler_2017**

The dot product of this matrix and the density vector, $\bar{\rho}$, gives the change in in the density vector. This is repeated over many small time periods to achieve the recorded results.



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

NISE (Numerical Integration of the Schrödinger Equation) is an existing open-source implementation of the simulation for these kinds of spectra.**nise** It was written by Kohler and Thompson while preparing their manuscript.**Kohler_2017**

NISE uses a slight variation on the algorithm presented, which allows for a 7-element state vector, but requires two simulations.

The end result is the same.

NISE is included as a reference for prior implementations.



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

- ▶ Use single 9×9 matrix rather than two 7×7 matrices.
- ▶ 99.5% of time is spent in highly parallelizable loop.
- ▶ 1/3 of time is spent in a single function, `ix_`.
 - ▶ Removed entirely, in favor of a simpler communication of what to record.
- ▶ Significant time in `rotor` function which computes $\cos(\theta) + i * \sin(\theta)$.
 - ▶ Replaced with $\exp(i * \theta)$, equivalent, more efficient, removed function call.
- ▶ Use variables to store and reuse redundant computations.

Resulted in almost an order of magnitude speed-up from algorithmic improvements alone. Remained highly parallelizable



Goals


Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations



Python cProfile trace, Single Core implementation, visualized with
SnakeViz.snakeviz



Goals


Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations



Python cProfile trace, Single Core implementation, visualized with
SnakeViz **snakeviz**



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

- ▶ NISE already had CPU multiprocessed parallelism, using Python standard library interfaces.
 - ▶ WrightSim inherited this CPU parallel implementation
 - ▶ Results in a 4x speed-up on a 4-core machine, almost no reduction due to Amdahl's law.
- ▶ A new Nvidia CUDA **Nickolls_2008** implementation.
 - ▶ Uses PyCUDA to call the kernel from within Python.
 - ▶ Just-in-time compiled (using nvcc) from C source code stored in Python strings.
 - ▶ Implementation differs slightly from pure Python implementation.
 - ▶ Only actively used Hamiltonians are held in memory, Python implementation computes all timesteps ahead of time.
 - ▶ Similarly, only the actively used electric fields are held in memory.
 - ▶ Hand written dot product and vector addition, rather than the numpy implementations.



Goals


Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations



Goals

Theory

NISE

Algorithmic
ImprovementsParallel
Implementations

Scaling Analysis

Limitations

- ▶ For low number of points, the CUDA implementation is of limited use.
 - ▶ Around 200 ms required for compilation.
 - ▶ 4-Core multiprocessed becomes faster below approximately 256 points.
 - ▶ CUDA implementation currently uses a hard coded block size of 256.
 - ▶ Only multiples of 256 may be used at present to avoid illegal memory access.
- ▶ Independant CUDA simulations are memory limited.
 - ▶ only a certain amount of memory can be allocated for a single CUDA process.
 - ▶ Each point in the simulation requires 500 complex numbers (represented as doubles) to be allocated
 - ▶ Additional data is needed, but dominated by this array.
 - ▶ This array must be transferred back to the host.
 - ▶ The limit is between 2^{18} and 2^{19} points

